

# CUDA Application Reduction Algorithm



Bedřich Beneš, Ph.D.  
Purdue University  
Department of Computer Graphics



## A Reduction Algorithm

- a reduction algorithm extracts a value from an array
- examples: sum, min, max, average...
- Let's assume data is in `int a[n];`

© Bedřich Beneš



## Sequential Reduction Algorithm

A sequential version  $O(n)$


```
x=action(a[0]);  
for (int i=1;i<n;i++)  
    x=action(a[i]);  
//now x stores the result
```

© Bedřich Beneš



## Parallel Reduction Algorithm

A parallel version (soccer tournament)  $O(\log_2 n)$



© Bedřich Beneš



## Parallel Reduction Algorithm

### Problems:

- partial results need to be shared
- very low arithmetic intensity

### Ideally:

- global synchronization

### CUDA:

- decompose into multiple kernel invocations



## Parallel Reduction Algorithm

### Metrics of performance:

- a) GFLOPs
- b) bandwidth

as the arithmetic intensity is low,  
bandwidth is a better measure



## Parallel Reduction Algorithm

- The operation will be performed per block
- Each block will load the data into shared memory
- All threads work on the block in parallel



## Comparison

- 32 MB of data ( $2^{22}$  integers)
- Tesla C1060
- Data block size 128
- Peak bandwidth 102 GB/sec

version \ values	Speedup	Time	Bandwidth

## Reduction Algorithm v.1

```
//all elements are in the shared memory
unsigned int n=2<<21;//# of integers
size_t size = n*sizeof(int);//datasize [B]
int maxThreads=128; //threads per block
int threads=(n<maxThreads)?n:maxThreads;
int blocks=n/threads; //# of blocks!!
dim3 dimBlock(threads,1,1);
dim3 dimGrid(blocks,1,1);
int smSize=threads*sizeof(int);//shared mem
//mallocs should be here
//note - size of the data is the block size
PR<<<dimGrid,dimBlock,smSize>>>(d);
```

© Bedrich Benes


## Reduction Algorithm v.1

```
__shared__ int sm[];
__device__ void PR(int *d){
    uint tid=threadIdx.x;
    uint i=blockIdx.x*blockDim.x+threadIdx.x;
    sm[tid]=d[i]; //copy to SM
    for (int stride=1;stride<blockDim.x;stride*=2)
    {
        __syncthreads();
        if (t%(2*stride)==0) sm[t]+=sm[t+stride];
    }
    if (tid==0) d[blockIdx.x]=sm[0]; //copy back
    //d[blockIdx.x] contains the sum of the block
}
```

© Bedrich Benes

## Reduction Algorithm v.1

- Single block parallel reduction




© Bedrich Benes

## Reduction Algorithm v.1

### Post process:

- The GPU kernel calculates data per *block*
- Results will be located in the first *block* elements of the global memory
- We will sum them with the same kernel...



© Bedrich Benes



## Reduction Algorithm v.1

```
if (blocks>1) toDo=1+blocks/128;
    else toDo=0;
for (int i=0;i<toDo;i++){
    threads=
        (blocks<maxThreads)?blocks:maxThreads;
    blocks=blocks/threads;
    dim3 dimBlock(threads,1,1);
    dim3 dimGrid(blocks,1,1);
    PR<<<dimGrid,dimBlock,smemSize>>>(d);
}
cudaMemcpy(hOut,d,s,cudaMemcpyDeviceToHost);
//hOut[0] is the result
```



## Comparison

- 32 MB of data ( $2^{22}$  integers)
- Tesla C1060
- Data block size 128
- Peak bandwidth 102 GB/sec

version \ values	Speedup	Time	Bandwidth
Version 1	1	7.800ms	4.28 GB/sec



## Reduction Algorithm v.1

Problems:

- thread divergence due to interleaved branch decisions  
if (t%(2\*stride)==0) sm[t]+=sm[t+stride];
- half of the threads does nothing!
- loop is expensive
- the % operator is slow



## Reduction Algorithm v.2

- replace the divergent branch with a non-divergent one


```
if ((tid%(2*stride))==0)
    sm[tid]+=sm[tid+stride];
```



```
int index=2*stride*tid;
if (index<blockDim.x)
    sm[index]+=sm[index+stride];
```

## Reduction Algorithm v.2

- Single block parallel reduction



© Bedrich Benes

## Comparison

- 32 MB of data ( $2^{22}$  integers)
- Tesla C1060
- Data block size 128
- Peak bandwidth 102 GB/sec

version \ values	Speedup	Time	Bandwidth
Version 1	1	7.800ms	4.28 GB/sec
Version 2	1.96	3.975ms	8.41 GB/sec

© Bedrich Benes

## Reduction Algorithm v.2

Problems:

- ~~thread divergence~~
- half of the threads does nothing!
- the % operator is slow
- shared memory bank conflicts
- loop is expensive

© Bedrich Benes

## Reduction Algorithm v.3

- replace the strided loop with a reversed one

```
for (stride=1; stride<blockDim.x; stride *= 2) {
    int ind=2*stride*tid;
    if (ind<blockDim.x) sm[ind]+=sm[ind+stride];
    __syncthreads();
}
```




```
for (stride=blockDim.x/2; stride>0; stride>>=1) {
    if (tid<stride) sm[tid]+=sm[tid+stride];
    __syncthreads();
}
```

© Bedrich Benes

## Reduction Algorithm v.3

- Single block parallel reduction



© Bedrich Benes

## Comparison

- 32 MB of data ( $2^{22}$  integers)
- Tesla C1060
- Data block size 128
- Peak bandwidth 102 GB/sec

version \ values	Speedup	Time	Bandwidth
Version 1	1	7.800 ms	4.28 GB/sec
Version 2	1.96	3.975 ms	8.41 GB/sec
Version 3	2.94	2.650 ms	12.43 GB/sec

© Bedrich Benes

## Reduction Algorithm v.3

Problems:

- ~~thread divergence~~
- half of the threads does nothing!
- ~~the % operator is slow~~
- ~~sm bank conflicts~~
- loop is expensive

© Bedrich Benes

## Reduction Algorithm v.4

let's make busy all threads in the first step

- 1) use only the half the blocks and
- 2) do the first reduction during the load from GM (replace single load with two loads)

© Bedrich Benes



## Reduction Algorithm v.4

```
uint tid=threadIdx.x;
uint i=blockIdx.x*(blockDim.x*2)+threadIdx.x;
sm[tid]= d[i]+d[i+blockDim.x];
__syncthreads();
for (stride=blockDim.x/2;stride>0;stride>>=1){
    if (tid<stride) sm[tid]+=sm[tid+stride];
    __syncthreads();
}
// write result for this block to global mem
if (tid == 0) d[blockIdx.x]=sm[0];
```

© Bedrich Benes



## Comparison

- 32 MB of data ( $2^{22}$  integers)
- Tesla C1060
- Data block size 128
- Peak bandwidth 102 GB/sec

version \ values	Speedup	Time	Bandwidth
Version 1	1	7.800 ms	4.28 GB/sec
Version 2	1.96	3.975 ms	8.41 GB/sec
Version 3	2.94	2.650 ms	12.43 GB/sec
Version 4	5.55	1.405 ms	23.90 GB/sec

© Bedrich Benes



## Reduction Algorithm v.4

- Problems
- It has low arithmetic intensity, so the bandwidth should be better
- But 23.9 GB/sec is about 24% of peak performance

© Bedrich Benes



## Reduction Algorithm v.5

- Let's unroll the loops!
- Number of active threads decreases with the number of iterations
- for stride $\leq$ 32 we have only one warp
- warp runs the same instruction (SIMD)
- for stride $\leq$ 32 \_\_syncthreads() is not necessary
- Let's unroll last 6 iterations

© Bedrich Benes



## Reduction Algorithm v.5

```
uint tid=threadIdx.x;
uint i=blockIdx.x*(blockDim.x*2)+threadIdx.x;
sm[tid] = d[i]+d[i+blockDim.x];
__syncthreads();
for (stride=blockDim.x/2;stride>32;stride>>=1)
{
    if (tid<stride) sm[tid]+=sm[tid+stride];
    __syncthreads();
}
```

© Bedrich Benes



## Reduction Algorithm v.5

```
if (tid < 32)
{
    sm[tid]+=sm[tid+32];
    sm[tid]+=sm[tid+16];
    sm[tid]+=sm[tid+8];
    sm[tid]+=sm[tid+4];
    sm[tid]+=sm[tid+2];
    sm[tid]+=sm[tid+1];
}
// write result for this block to global mem
if (tid == 0) d[blockIdx.x]=sm[0];
}
```

© Bedrich Benes



## Comparison

- 32 MB of data ( $2^{22}$  integers)
- Tesla C1060
- Data block size 128
- Peak bandwidth 102 GB/sec

version \ values	Speedup	Time	Bandwidth
Version 1	1	7.800 ms	4.28 GB/sec
Version 2	1.96	3.975 ms	8.41 GB/sec
Version 3	2.94	2.650 ms	12.43 GB/sec
Version 4	5.55	1.405 ms	23.90 GB/sec
Version 5	5.89	1.325	24.33 GB/sec

© Bedrich Benes



## Additional measurements

- 128 MB of data ( $2^{24}$  integers)
- Tesla C1060
- Data block size 512
- Bandwidth 50 GB/sec
- That is 50% of peak bandwidth 102 GB/sec

© Bedrich Benes





## Conclusions

- Optimizations lead to a great speedup
- Algorithmic optimizations (addressing, cascading, etc) ~ 3x speedup
- Loop unrolling another 2x speedup
- First optimize, then unroll the loops!



## Reading

- Harris, M., *Optimizing Parallel Reduction in CUDA* (NVIDIA Dev. technology)
- Kirk, D.B., Hwu, W.W., *Programming Massively Parallel Processors*, NVIDIA, Morgan Kaufmann 2010

